

ANVCS — AI-Native Version Control System (Core Spec v0.4)

Subtitle: Offline-First, Disconnected CRDT Mode, Secure Remote Connectivity, and Enterprise Resilience

This v0.4 release consolidates the core ANVCS design and **integrates the new CRDT-disconnected (offline) mode**, plus CLI/UX flows, security, and testability guarantees for air-gapped environments. It is intended to replace earlier versions as the **main document**.

0) Executive Summary

ANVCS is a **Git-compatible, AI-native version control system** that combines **CRDT live collaboration**, a **commit DAG**, **semantic merge**, and **rich AI provenance** (prompts, model versions, tests, artifacts). New in v0.4: a formal **offline/CRDT-disconnected mode** that maintains **full parity** with connected workflows, ensuring developers and agents can work **without network access**, commit locally, and **rejoin** later without losing context or integrity.

Core value - **Git immutability + compatibility** preserved (content-addressed commits; push/pull; CI unaffected) - **CRDT live edits** across agents with conflict-free convergence - **Semantic merges** at commit boundaries (AST/test/LLM-assisted) - **Provenance** captured per commit (prompt, model, tests, artifacts) - **Offline-first parity**: connected and disconnected experiences behave **the same** - **Pluggable artifact storage** (S3 versioning/IPFS/NFS), exact version pinning

1) Problem Statement (recap)

- Git was optimized for human-paced teams, not 24/7 swarms of agents. Merge conflicts, history opacity, and cross-repo drift become severe.
 - Non-text assets (models, datasets) lack reproducibility without external pinning.
 - Enterprises need **offline/air-gapped** workflows while retaining auditability and consistency.
-

2) Goals & Non-Goals

Goals - Preserve Git **immutability**, **content addressability**, and **interop**. - Enable **multi-agent** concurrency with low friction. - Make **intent traceable**: prompts, models, test outcomes per commit. - Support **any file type** via **version-aware storage adapters** (S3 VersionId, IPFS CID, etc.). - Provide **offline-first parity**: commits and metadata identical on/offline. - Secure, observable, testable in **air-gapped** and **cloud** settings.

Non-Goals (v0.4) - Replacing Git hosting; blockchain consensus layers; automatic resolution of all semantic conflicts.

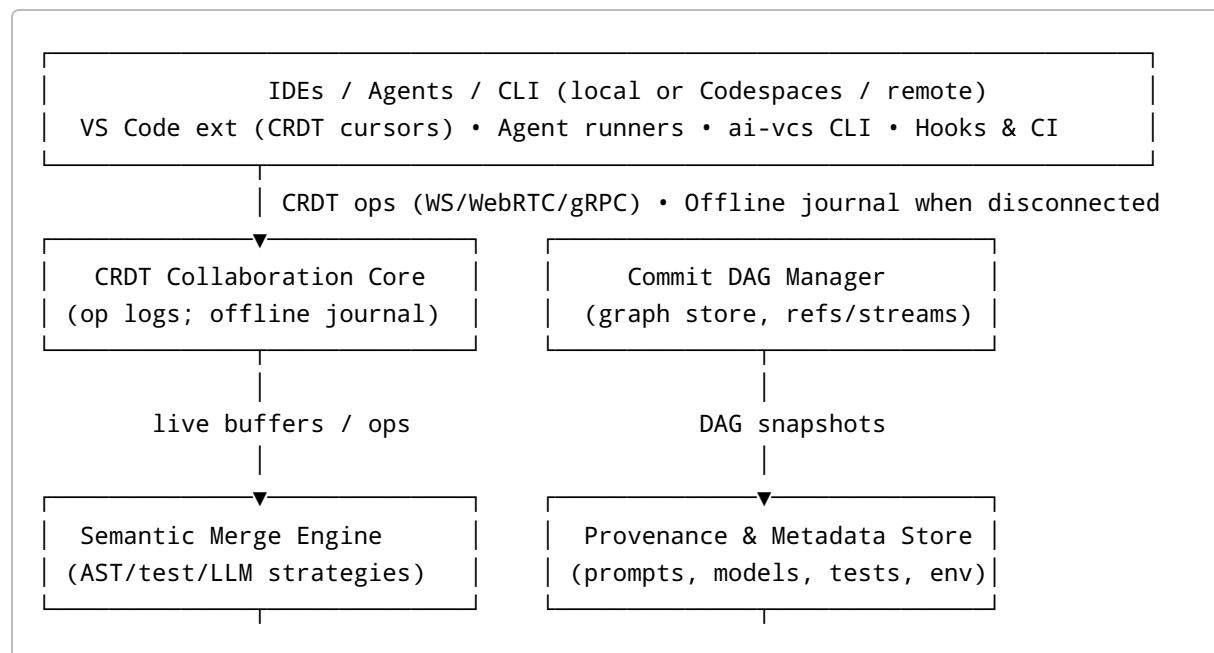
3) Core Principles

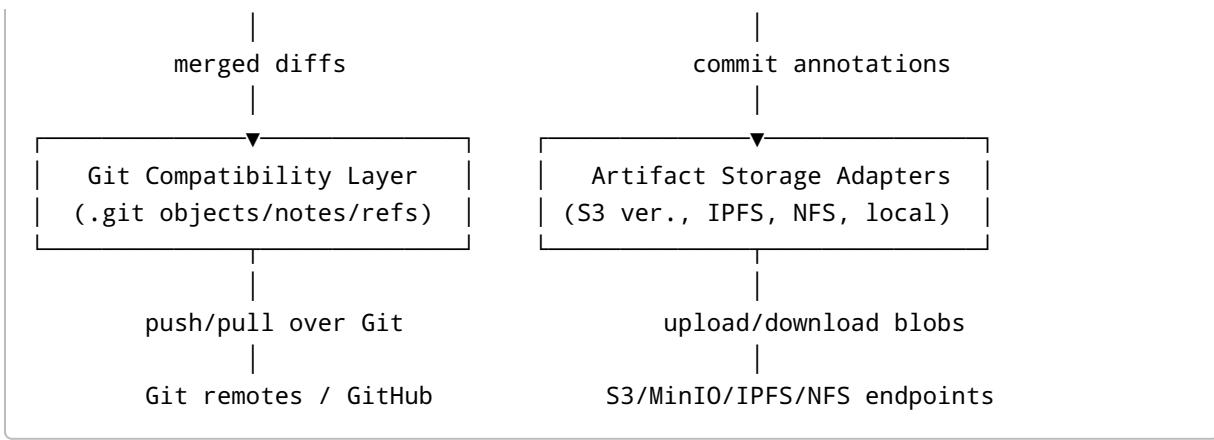
1. **Git-Compatibility First** — Dual-write to `.git/` on every commit; Git notes/structured messages for AI metadata.
 2. **Immutability & Reproducibility** — Append-only DAG; content-hash; pinned artifacts `{key, version_id, sha256}`.
 3. **CRDT for Edits, DAG for Versions** — Ops converge live; commits snapshot coherent states.
 4. **Provenance as First-Class** — Capture prompt/model/tests/env/artifacts for every commit.
 5. **Semantic Merge** — AST/test-guided merges at DAG boundaries.
 6. **Offline-First Equivalence** — Connected and disconnected modes **produce identical artifacts and metadata**.
 7. **Pluggable Everything** — Storage, transports, merge strategies, policies, identity.
-

4) Conceptual Model

- **CRDT Room** — Operation-based per-file logs; edits stream among participants when connected.
 - **Commit (Semantic Snapshot)** — Agents/devs take snapshots at task boundaries; creates DAG node + Git commit with AI metadata.
 - **DAG** — Directed acyclic graph with multi-parent merges; branchesstreams are just **named pointers**.
 - **Agent Context** — Structured metadata: prompt, model, tests, coverage, env, artifacts.
 - **Artifact Store** — Offloads large binaries via adapters (S3/IPFS/NFS) with exact version pinning.
-

5) Architecture Overview





6) Data Model

Commit object (concept)

```

commit_id: <hash>
parents: [<hash>, ...]
author: { name, email, time }
message: |
    Summary line

ai-meta:
    prompt: "..."
    model: "..."
    tests_passed: true
    test_report_ref: ".ai/meta/<id>/tests.json"
    coverage: 87.1
    env: { python: "3.12", node: "24", os: "ubuntu-24.04" }
    agent_id: "agent-foo"
    artifacts:
        - type: model
          storage: s3
          bucket: anvcs-artifacts
          region: eu-west-1
          key: models/agent_v1.pt
          version_id: 3HL4IxNZA...
          sha256: 4f34...9a1c
          size: 104857600

```

CRDT ops (per file) — operation journal (insert/delete/replace) with lamport/vector clocks and stable op IDs; compacted over time.

DAG store — append-only nodes + edges; refs map; cross-repo meta-commit manifests (optional).

7) CRDT-Disconnected (Offline) Mode — Local-First & Air-Gapped Resilience

Design objective: When CRDT connectivity is unavailable, **nothing breaks**. Edits apply to a **local CRDT journal**, and every `ai-vcs commit` produces the **same** DAG node + Git commit + AI metadata as if online.

Key guarantees - Parity: Connected vs disconnected commits are **identical in structure** (Git commit + `.ai/meta` schema + artifact pins). - **Dual-write:** Each commit always writes to `.git` (objects/refs) and `.ai/` locally. No daemon required. - **No special commands:** Continue using `ai-vcs commit`. If disconnected, networking is skipped; commit completes locally. - **Rejoin later:** `ai-vcs rejoin-room` (or reconnect) publishes buffered ops/commits, merges DAGs, hydrates peers. - **Air-gapped ready:** Works fully offline with local storage (MinIO/NFS) and private Git remotes; CI can run on isolated networks.

Workflow 1) Developer/agent edits files → ops appended to **local CRDT journal**. 2) `ai-vcs commit` → snapshot current state → **DAG node + Git commit + AI metadata** are created locally (identical schema). 3) Artifacts matched by `blobs.track` are stored via **local adapter** (e.g., MinIO) and pinned by `{key, version_id, sha256}`. 4) On reconnect: publish ops; pull peers' ops; CRDT ensures convergence; semantic merge resolves overlapping edits at DAG boundaries.

Visibility & merge when rejoining - Offline branches appear as separate DAG heads; semantic merge can combine them. - All provenance (prompts/tests) is preserved for review.

Testability - Provide offline test scenarios: disable network; ensure `commit`, `log`, `graph`, `checkout` work; artifact hydration uses only local endpoints; rejoin test validates convergence and merge.

8) Semantic Merge (recap)

- **AST-aware structured diff/merge** (renames/moves/extractions understood);
 - **Guards** with typechecks/tests/linters;
 - **LLM-assisted proposals** for tricky overlaps (gated by tests);
 - **Non-code** drivers for JSON/YAML, notebooks, IDLs.
-

9) Storage Adapters (Any File Type)

- **Interface:** `upload() → {key, version_id, sha256, size}`,
`download(key, version_id)`, `exists()`, `delete_version()`.
- **S3 Versioning:** Stable key; commit pins `VersionId`; exact retrieval on checkout; supports MinIO/Spaces.

- **Local/Hybrid modes:** Cache under `.ai/blobs/`; defer sync; resilient retries; proxy/env awareness.

Config (`ai-vcs.config.yaml`)

```
storage: { mode: hybrid, primary: s3, fallback: local }
adapters:
  s3:
    bucket: ${AI_VCS_S3_BUCKET}
    region: ${AWS_REGION}
    endpoint_url: ${S3_ENDPOINT_URL:-}
    use_presigned: true
  blobs:
    track: ["*.pt", "*.onnx", "datasets/**/*.{csv,ipynb,png}", "**/*.mp4"]
    upload_on_commit: true
    compress: true
```

10) Git Interop Strategy

- **Dual-write** every commit to Git (tree/parents) + `.ai/meta` (or Git notes);
- **Refs** map DAG heads to Git branches;
- **Push/Pull** use normal Git remotes; metadata rides via notes or repo files.

11) Security, Identity, Governance

- **Commit signing** (Ed25519/GPG); per-agent keypairs; short-lived tokens; OIDC/STS for storage.
- **Scopes:** `read|write|commit|merge`; path allow/deny policies.
- **Audit:** immutable logs of snapshot/merge decisions; provenance captured per commit.

12) CLI, API & VS Code UX

CLI (selected)

```
ai-vcs init
ai-vcs collab start|join <room>      # optional; not required to commit
ai-vcs add <paths>
ai-vcs commit -m "feat: X"
  --prompt "Implement X" --model gpt-4o --test-report ./ai/meta/tests.json
ai-vcs merge <src> into <dst> [--auto --test-check]
ai-vcs graph --since 2w --filter agent:foo
```

```

ai-vcs rejoin-room          # publish buffered ops/commits when back
online
ai-vcs push|pull origin main

```

VS Code extension - Live cursors & diff heatmap (when connected) - **Preview → Confirm** for snapshots (semantic summary, risk flags) - **Playback timeline** (op journal), **DAG view**, **Artifacts pane** - **Offline badge** with graceful degradation (local only)

Agent Streaming API (summary) - WebSocket/WebRTC/gRPC ops; **snapshot** to persist commit; **artifact** announcements; acks/backpressure; auth scopes.

13) Deployment Topologies (Secure & Flexible)

| Topology | When | Pros | Cons |
|--------------------------------|-------------------------|---------------------------|------------------------|
| Local-only | Solo/offline | Zero infra, fully offline | No live collab |
| Ephemeral relay (devcontainer) | Small teams/ Codespaces | Easy, low-latency | Process lifecycle mgmt |
| WebRTC P2P | Cross-site agents | E2E, low relay cost | NAT/TURN complexity |
| Managed relay cluster | Enterprises | SSO, quotas, audit | Operate service |

Connectivity: TLS/DTLS/mTLS; JWT/OIDC/SSH-sig; ICE (STUN/TURN); proxy envs respected.

14) Observability & Debuggability

- **Metrics:** ops/sec, latency, queue depth, merge success, artifact hydration.
- **Session journal:** join/leave, batch IDs, snapshot/merge attempts, test outcomes.
- **Time-travel replay:** rebuild worktree at op index/timestamp.
- **Tombstone GC** for CRDT stores; bounded logs.

15) Multi-Repo & Cross-Modality (optional)

- **Meta-commit groups:** sign/pin coordinated commits across repos; CI validates group.
- **Modality-aware diffs:** JSON/YAML (schema), notebooks (cell graph), IDLs (symbols), media manifests.

16) Test Matrix (incl. Air-Gapped)

- **Offline unit/integration:** no network; ensure `commit/log/graph/checkout` pass; artifacts served by local endpoints only.
- **Rejoin convergence:** create divergent offline branches; rejoin; assert CRDT convergence + semantic merge correctness.
- **Codespaces + proxy:** ensure adapter respects `HTTPS_PROXY/NO_PROXY`.
- **S3 versioning:** verify `VersionId` pinning, exact hydration on checkout.

17) Roadmap

- v0.5: cross-repo meta-commit manifests; policy DSL; extended modality drivers.
- v0.6: P2P hardening; delta-sync optimizations; fine-grained blame w/ provenance UI.

18) Appendix — Interfaces & Schemas (extract)

BlobStore (version-aware)

```
class BlobStore(ABC):  
    def upload(local_path: str, key: str) -> Dict: ...  
    def download(key: str, version_id: Optional[str], dest_path: str) ->  
        None: ...  
    def exists(key: str, version_id: Optional[str] = None) -> bool: ...  
    def delete_version(key: str, version_id: str) -> None: ...
```

S3VersionedBlobStore — returns `{key, version_id, sha256, size}`; supports MinIO/Spaces via `endpoint_url`.

Agent Streaming (WS) — envelope

```
{ "type": "op|ack|snapshot|artifact|presence|heartbeat|error", "room": "...",  
"agent": "...", "seq": 512, "data": { ... } }
```

Commit metadata (JSON)

```
{  
    "commit_id": "abc123...",  
    "ai_meta": {  
        "prompt": "Implement validation",  
        "model": "gpt-4o-2025-10",
```

```
    "tests_passed": true,
    "coverage": 91.2,
    "agent_id": "agent-validate",
    "artifacts": [
        {"storage": "s3", "bucket": "anvcs-artifacts", "region": "eu-west-1",
         "key": "models/validator.pt", "version_id": "3HL4IxNZ...", "sha256":
         "4f34...9a1c", "size": 104857600 }
    ]
}
```

Summary — v0.4 elevates ANVCS to a **resilient, offline-first** system: identical commit semantics on/offline, secure remote connectivity, and enterprise-grade observability — all while retaining Git compatibility and adding AI-native collaboration.