



Git Limitations in Multi-Agent Workflows

- **Opaque history (reflog & local state):** Git's reflog and stashes are local-only and not shared across collaborators, making it hard to audit automated agent actions or recover context. Hidden operations (like resets or reverts) disappear from the shared history, leaving gaps in traceability.
- **Excess commit churn:** AI agents can generate *very* high commit volumes. For example, one reported multi-agent pipeline made *over 3,000 commits in 48 hours*, rendering the repository “unrecognizable” ¹. Such churn floods the history with trivial changes, obscuring human-readable milestones.
- **Merge/cherry-pick fragility:** Simultaneous edits by agents lead to frequent merge conflicts. Traditional Git merging or cherry-picking is brittle with automated changes, often requiring manual resolution. In contrast, CRDT-based systems (like CodeCRDT) achieve *strong eventual consistency* with “0% character-level conflicts” and no manual merges ². Git’s snapshot/branch model was not designed for hundreds of concurrent contributors and thus breaks down under agent swarm.
- **Lack of contextual blame:** Standard `git blame` only attributes lines to an author name, without any metadata about AI involvement. When AI writes or transforms code, blame loses meaning. As one author notes, “That line might’ve been written by GPT-4... no one’s quite sure who’s on the hook” ³. Important context (e.g. the prompt, model version, or reviewing agent) is missing, making accountability and debugging very hard.

Extending the CRDT+DAG Model

To address these Git shortcomings, ANVCS should enrich its CRDT-based collaboration model with a richer version graph and metadata:

- **Graph-based version history:** Instead of a linear commit log, maintain a *full directed acyclic graph* (DAG) of versions. Each agent’s edit produces a new node in the graph with edges denoting “fork” or “merge” operations. This is the approach of EvoGit, which tracks *all* versions as graph nodes so agents can “asynchronously read from and write to the evolving code repository” ⁴. Fine-grained branching and implicit concurrency allow multiple agents to work in parallel without blocking one another ⁴. The DAG preserves every edit’s lineage, preventing the loss of context that happens when rebasing or amending in Git.
- **Atomic operation commits:** Capture CRDT edits as atomic operations or deltas rather than raw file diffs. For example, log each CRDT operation (insertion, delete, etc.) with causality information. This “operation journal” can be periodically snapshotted into Git commits, but the underlying units remain CRDT-operations. This avoids large monolithic patches and makes conflict resolution trivial. Systems like DeltaDB and CodeCRDT use operation-based syncing to “eliminate conflicts entirely through automatic CRDT resolution” ⁵ ².
- **Contextual metadata & Git notes:** Attach rich metadata to each version node/commit. Using Git’s extensibility (e.g. [Git notes](#)), ANVCS can record the prompt text, model version, responsible agent ID, execution results, and test diagnostics alongside the commit. As one proposal notes, EvoGit “records all changes as immutable commits” and “enhances traceability by attaching rich metadata (e.g. build diagnostics, agent identifiers) via Git notes” ⁶. This means a commit isn’t just a diff, but a bundle of

context for why and how it happened. Agents could even embed pointer references to parts of the CRDT log that produced the change, creating a continuous chain of provenance.

- **Temporal checkpoints:** Incorporate lightweight checkpoints or snapshots for recovery and blame. For example, emulate a Git-Context-Controller (GCC) style system where an agent's working memory is versioned in the repo ⁷. GCC "creates checkpoint systems for context retrieval, enabling agents to maintain conversation history and decision context linked to code evolution" ⁷. ANVCS could similarly treat each agent's local state (prompts, variables, intermediate results) as versioned files or branches, ensuring that even AI "thought processes" are auditable.

By combining CRDT merging at the file level with a semantically rich commit DAG, ANVCS can mitigate Git's linear-log limitations. The DAG makes parallel development explicit (no hidden reflogs), and metadata transforms opaque blame into "who changed what, why, and under which AI policy" for complete accountability ⁶ ³.

Real-Time Human Visibility

Human developers must see exactly what agents do *before* it lands in the main codebase. ANVCS can add GUIs and workflows such as:

- **Live diff review:** Provide a multi-agent diff interface. For example, Claude Code proposes a *parallel agents* mode where each agent runs in its own Git worktree, and once complete the UI presents a **side-by-side multi-diff**: the original code on one pane and each agent's proposed changes in other panes ⁸. The developer can then choose to apply one agent's solution wholesale, or cherry-pick pieces. Warp Terminal similarly highlights the importance of a "live diff view" that lets the developer "steer agents in-flight," turning the AI from a black box into a transparent collaborator ⁹. ANVCS should integrate a diff-preview window where every pending CRDT commit or branch tip can be inspected and approved.
- **Commit preview mode:** Before finalizing a commit, show a preview of its effects. This two-step "Preview → Confirm" workflow (advocated in safety-oriented AI frameworks) lets the user see exactly *what will happen*, who will be affected, and how to undo it. (Some have suggested a hold-to-confirm mechanism before autonomous commits.) ANVCS could display the incremental diff or a summary log of operations and require explicit human confirmation on significant changes, preventing blind or damaging updates.
- **Live collaborative editor:** Embed a real-time collaborative code editor (e.g. VS Code/Monaco) wired to the CRDT. Showing agents' cursors and highlights in real-time (as in Fig.2 of CodeCRDT) allows developers to watch agents code as it happens ¹⁰. In CodeCRDT's demo, colored cursors indicate multiple agents editing concurrently while Yjs CRDT syncs everyone's view ¹⁰. ANVCS could adopt a similar UI: each agent's edits appear in color or with tags. A human could even jump in to edit alongside agents, or pause/resume an agent mid-edit.
- **Change playback/timeline:** Record the sequence of CRDT edits over time, enabling a "replay" mode. This could be implemented by logging timestamped CRDT operations or commit events. A UI timeline would let developers scrub backward/forward through the code evolution history, viewing how the document grew. Screenshots of the IDE or an event log (e.g. "Agent A inserted function foo() at t=10:01") would support forensic analysis.
- **Agent activity dashboard:** Show a timeline or chart of agent actions and presence. For instance, a panel listing active agents with status, steps performed, and last update timestamp (as Claude Code suggests ¹¹). Visual indicators (✓ or ✗) for successful commits, merge attempts, or test results can

alert the user to problems. Integration with a chat or feed (e.g. Slack-like) could notify the dev when an agent finishes a task or requires attention.

Collectively, these features turn the developer's role from passive reviewer into active supervisor. They align with industry proposals for "streaming" AI changes rather than opaque bulk commits ⁹ ⁸.

Remote Agent Connectivity

ANVCS must work when agents run off-site or in the cloud. We recommend peer-to-peer and relay-based options instead of a permanent central server:

- **WebRTC / P2P sync:** Use WebRTC (with STUN/TURN) for direct, encrypted CRDT sync between the dev's machine and remote agents. This is essentially what the open-source Conclave editor does: "*CRDT and WebRTC based real-time, peer-to-peer, collaborative text editor*" ¹². Similarly, each ANVCS node (developer or agent) could be a peer in a mesh, syncing via Yjs or Automerge over WebRTC. A lightweight signaling broker would establish the initial connection, but thereafter updates flow directly, with NAT traversal handled by ICE. This eliminates the need for a continuously-running CRDT server.
- **On-demand relay/proxy:** If direct P2P isn't feasible (e.g. strict firewalls), run a temporary relay. For example, the local dev machine could expose a TLS-secured port (via SSH tunneling, ngrok, or a Kubernetes-inlet) and the agent can connect to it. Alternatively, use a TURN server (common in WebRTC) to forward traffic while keeping end-to-end encryption. The key is that the relay only needs to exist during the session; ANVCS connections can tear down when idle.
- **Overlay network / VPN:** In corporate environments, agents could join the developer's private network via VPN or libp2p/mesh protocols. Each agent would get a stable address, allowing the ANVCS nodes to discover each other directly. Technologies like WireGuard tunnels or peer-to-peer networks (e.g. IPFS/libp2p) can facilitate this.
- **Mobile CRDT servers:** As a hybrid, a dev might run a lightweight CRDT "server" on their machine only while collaborating. Remote agents push updates to this server process, which then merges and persists them. When agents disconnect, the server can dump any remaining ops to local storage. This isn't a full-time centralized server (it runs only when needed) but ensures reliable sync behind corporate firewalls.

In all cases, traffic must be secured (DTLS/TLS, authenticated via tokens or keys). By leveraging P2P CRDT libraries and modern NAT traversal, ANVCS can support remote agents without a constant central backend

¹².

Resilience, Observability and Debuggability

Finally, ANVCS should be engineered like a distributed system with robust logging and fail-safes:

- **Eventual consistency & durability:** Rely on CRDT's fault tolerance: even if agents or the network go offline, all updates are eventually applied without conflict. As CodeCRDT demonstrates, "**strong eventual consistency**" means "all agents converge to the same state with zero merge failures" ². ANVCS nodes should persist CRDT state locally and replay operations on restart. Automatic garbage collection of tombstones (in CRDTs like Yjs) ensures unbounded history doesn't bloat storage.

- **Immutable audit log:** Keep the full commit/DAG history as an immutable record. EvoGit emphasizes that “*all changes are recorded as immutable commits*” enabling robust traceability ¹³. ANVCS should do likewise: every DAG node is write-once. Using Git’s model, each merge or release to “main” is a permanent, signed reference to a CRDT snapshot. This allows unambiguous rollbacks and post-mortems.
- **Rich metadata logging:** Record **everything** into the log or monitoring system. Based on EvoGit’s traceability appendix, ANVCS should attach diagnostics (build/test results, performance metrics, agent IDs) to each commit ⁶. Logging the CRDT operation ID and timestamp for each change provides a replayable history. In practice, that could mean writing structured JSON logs or using Git notes for each commit to store contextual metadata (as EvoGit does).
- **Context checkpoints:** Periodically checkpoint the agent’s state and prompts into the repo. As with Git-Context-Controller, a serialized “memory snapshot” (openAI chat history, local variables, etc.) can be versioned alongside code ⁷. This lets us later inspect exactly what the agent “knew” when it made a change. A human debugging an issue could check out an older “memory branch” to see the agent’s full context at that point.
- **Monitoring & alerts:** Instrument CRDT sync health, latency, and agent errors. For example, track metrics like “time since last CRDT sync,” “dropped messages,” or “merge conflicts escalated.” ANVCS could expose Prometheus metrics or logs that trigger alerts if an agent is unresponsive or if the DAG becomes disconnected. Visualization tools (e.g. a live commit-graph dashboard) help spot anomalies. EvoGit points out that such traceability “enables rigorous auditing” and even per-agent performance evaluation ¹⁴.
- **Debugging support:** Provide tools to inspect the live CRDT state. For instance, ANVCS could include a REPL or GUI command to dump the current document state, pending operations, or version history. Embedding trace IDs in log messages allows tracing a bug from test failure back through the chain of commits/ops that caused it. In tricky cases, one could “time-travel” by replaying CRDT events from a certain point to reproduce a misbehavior in isolation.

By applying distributed-systems best practices (append-only logs, monitoring, encryption) to the CRDT+DAG stack, ANVCS becomes self-describing and recoverable. Every code change is tied to metadata and visible to the human in near-real-time, preventing silent failures. This combination of an immutable Git-derived DAG ¹³ and CRDT synchronization ² ensures that agent-driven development remains both extensible *and* transparent.

References: Recent work on AI-native VCS and multi-agent development (EvoGit ⁴ ⁶, CodeCRDT ², Warp AI ⁹, etc.) informed these proposals. Those systems underline the value of DAG-based versioning with rich metadata for traceability and real-time collaborative editing with CRDTs to eliminate conflicts and boost resilience.

¹ When AI Tools Fight Each Other: The Hidden Chaos of Multi-Agent Workflows | by Tech Digest HQ | Aug, 2025 | Medium

<https://medium.com/@techdigesthq/when-ai-tools-fight-each-other-the-hidden-chaos-of-multi-agent-workflows-83169e8dcc6f>

² ¹⁰ 1 Introduction

<https://arxiv.org/html/2510.18893v1>

³ The New Git Blame: Who’s Responsible When AI Writes the Code? | by Alissa | PullFlow | Medium

<https://medium.com/pullflow/the-new-git-blame-whos-responsible-when-ai-writes-the-code-659b7ae60e71>

4 6 13 14 [2506.02049] EvoGit: Decentralized Code Evolution via Git-Based Multi-Agent Collaboration
<https://arxiv.labs.arxiv.org/html/2506.02049v1>

5 7 How Git Usage and DVCS Are Evolving in the AI Age with Next-Generation Version Control Systems - SoftwareSeni

<https://www.softwareseni.com/how-git-usage-and-dvcs-are-evolving-in-the-ai-age-with-next-generation-version-control-systems/>

8 11 [FEATURE] Parallel Multi-Agent Workflows for Code Generation and Planning · Issue #10599 · anthropics/clause-code · GitHub

<https://github.com/anthropics/clause-code/issues/10599>

9 Warp AI: My In-Depth Guide to the Agentic Terminal Changing the Future of Coding

<https://skywork.ai/skypage/en/Warp-AI:-My-In-Depth-Guide-to-the-Agentic-Terminal-Changing-the-Future-of-Coding/1974364910407839744>

12 GitHub - conclave-team/conclave: CRDT and WebRTC based real-time, peer-to-peer, collaborative text editor

<https://github.com/conclave-team/conclave>